

# Probabilistic Programming and AI: Lecture 2

## Introduction to Probabilistic Programming Languages

---

Markus Böck and Jürgen Cito

Research Unit of Software Engineering

# Table of contents

1. Recap on Bayesian Inference
2. Generative Modelling
3. Probabilistic Programming Languages
4. Implementing a minimal PPL
5. First General-Purpose Inference

## Recap on Bayesian Inference

---

## Recap on Bayesian Inference

- Latent variables  $\Theta$
- Observed variables  $X$
- Observes are believed to be generated from latents

# Recap on Bayesian Inference

- Latent variables  $\Theta$
- Observed variables  $X$
- Observes are believed to be generated from latents
- Prior distribution  $P(\Theta)$  captures our belief/information about the latent variables before observing data

# Recap on Bayesian Inference

- Latent variables  $\Theta$
- Observed variables  $X$
- Observes are believed to be generated from latents
- Prior distribution  $P(\Theta)$  captures our belief/information about the latent variables before observing data
- Likelihood  $P(X|\Theta)$  encodes how we believe the observes are generated from latents

# Recap on Bayesian Inference

- Latent variables  $\Theta$
- Observed variables  $X$
- Observes are believed to be generated from latents
- Prior distribution  $P(\Theta)$  captures our belief/information about the latent variables before observing data
- Likelihood  $P(X|\Theta)$  encodes how we believe the observes are generated from latents
- Posterior distribution is obtained with Bayes Theorem

$$P(\Theta|X) = \frac{P(X|\Theta) \times P(\Theta)}{P(X)}$$

# Recap on Bayesian Inference

- Latent variables  $\Theta$
- Observed variables  $X$
- Observes are believed to be generated from latents
- Prior distribution  $P(\Theta)$  captures our belief/information about the latent variables before observing data
- Likelihood  $P(X|\Theta)$  encodes how we believe the observes are generated from latents
- Posterior distribution is obtained with Bayes Theorem

$$P(\Theta|X) = \frac{P(X|\Theta) \times P(\Theta)}{P(X)}$$

- Posterior is the updated belief, our information about the latents after having observed data



# Recap on Bayesian Inference

Prior  $P(\Theta)$

- design decision: can be constructed from expert knowledge / previous experiments
- **computable**

# Recap on Bayesian Inference

Prior  $P(\Theta)$

- design decision: can be constructed from expert knowledge / previous experiments
- **computable**

Likelihood  $P(X|\Theta)$

- encodes generative process from latents  $\Theta$  to observes  $X$
- **computable**

# Recap on Bayesian Inference

## Prior $P(\Theta)$

- design decision: can be constructed from expert knowledge / previous experiments
- **computable**

## Likelihood $P(X|\Theta)$

- encodes generative process from latents  $\Theta$  to observes  $X$
- **computable**

## Joint Probability $P(X, \Theta)$

- product of prior and likelihood  $P(X, \Theta) = P(X|\Theta) \times P(\Theta)$
- **computable**

# Recap on Bayesian Inference

Marginal Likelihood / Model Evidence  $P(X)$

- can be interpreted as the probability of generating the observed data for all possible values of the latent variables,  
$$P(X) = \int P(X, \Theta) d\Theta,$$
- it can be understood as the probability of the model itself
- can only be **approximated** in most models

# Recap on Bayesian Inference

## Marginal Likelihood / Model Evidence $P(X)$

- can be interpreted as the probability of generating the observed data for all possible values of the latent variables,  
 $P(X) = \int P(X, \Theta) d\Theta,$
- it can be understood as the probability of the model itself
- can only be **approximated** in most models

## Posterior $P(\Theta|X)$

- is the updated belief, the information about the latents after having observed data
- **what we want to know**
- can only be **approximated** in most models

# Example

## Coin model

$$p \sim \text{Uniform}(0, 1)$$

*latent*

$$x_i \sim \text{Bernoulli}(p)$$

*observed*

# Example

## Coin model

$$p \sim \text{Uniform}(0, 1) \quad \text{latent}$$

$$x_i \sim \text{Bernoulli}(p) \quad \text{observed}$$

$$\Theta = p, X = (x_1, \dots, x_n), P(X, \Theta) = P(p)P(x_1|p) \cdots P(x_n|p)$$

# Example

## Coin model

$$p \sim \text{Uniform}(0, 1) \quad \text{latent}$$

$$x_i \sim \text{Bernoulli}(p) \quad \text{observed}$$

$$\Theta = p, X = (x_1, \dots, x_n), P(X, \Theta) = P(p)P(x_1|p) \cdots P(x_n|p)$$

$$\underbrace{\log P(X, \Theta)}_{\text{"log-prob"}} = \underbrace{\log P(p)}_{\text{log-prior}} + \underbrace{\sum_{i=1}^n \log P(x_i|p)}_{\text{log-likelihood}}$$



# Example

## Coin model

$p \sim \text{Uniform}(0, 1)$  *latent*

$x_i \sim \text{Bernoulli}(p)$  *observed*


$\Theta = p, X = (x_1, \dots, x_n), P(X, \Theta) = P(p)P(x_1|p) \cdots P(x_n|p)$

$$\underbrace{\log P(X, \Theta)}_{\text{"log-prob"}} = \underbrace{\log P(p)}_{\text{log-prior}} + \underbrace{\sum_{i=1}^n \log P(x_i|p)}_{\text{log-likelihood}}$$

```
1 def prior(p):
2     return dist.Uniform(0,1, validate_args=False).log_prob(p).exp()
3
4 def likelihood(xs, p):
5     lp = torch.tensor(0.)
6     for x in xs:
7         lp += dist.Bernoulli(p, validate_args=False).log_prob(x)
8     return lp.exp()
9
10 def joint(xs, p):
11     return prior(p) * likelihood(xs, p)
12
13 xs = torch.tensor([0., 1., 1., 0., 0.]
```

# Example

```
1 ps = torch.linspace(0,1,200)
2 posterior_unnormalised = torch.hstack([joint(xs, p) for p in ps])
3 marginal = torch.trapz(posterior_unnormalised, ps)
4 posterior = posterior_unnormalised / marginal
```



# Example

## Compared to Frequentist Statistics

Maximum-Likelihood-Estimator:

```
ps = torch.linspace(0,1,10000)
l = torch.hstack([likelihood(xs, p) for p in ps])
ps[l.argmax()]
```

**Returns:**

0.3999

# Example

## Compared to Frequentist Statistics

Maximum-Likelihood-Estimator:

```
ps = torch.linspace(0,1,10000)
l = torch.hstack([likelihood(xs, p) for p in ps])
ps[l.argmax()]
```

**Returns:**

0.3999

Hypothesis testing  $\mathcal{H}_0 : p = 0.5$ :

```
m = xs.mean() # test statistic: mean
means = dist.Bernoulli(0.5).sample((len(xs), 10000)).mean(dim=0)
a = min(1-m,m)
b = max(1-m,m)
# probability that sample is more extreme than observations
((means < a) | (b < means)).float().mean()
```

**Returns:**

tensor(0.3662) # p-value -> cannot reject null

# What is conditional probability

# What is conditional probability


It is a "restriction" of the sample space:

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

<https://seeing-theory.brown.edu/compound-probability/index.html#section3>

# What is conditional probability

## Prevalence of a disease



$$\#(\text{Sick} \mid \text{Total}) \equiv \#(\text{Sick}) \\ = 10$$

$$p(\text{Sick}) = 10/\#(\text{Total}) \\ = 10/100 \\ = 1/10$$


[https://medium.com/@javiergb\\_com/](https://medium.com/@javiergb_com/)

[why-testing-positive-for-a-disease-may-not-mean-you-are-sick-4a3a16a290eb](#)

# What is conditional probability


## Results of the test

Test Positive Test Negative






# What is conditional probability



$$\#(\text{Positive}) = 18$$


$$p(\text{Positive}) = 18/\#(\text{Total}) \\ = 18/100$$



$$p(\text{Sick} \mid \text{Positive}) = 9/18$$


How to calculate it? → BAYES THEOREM!

# What is conditional probability



#(Positive) = 18

$p(\text{Positive}) = 18/\#(\text{Total})$   
 $= 18/100$



$p(\text{Sick} | \text{Positive}) = 9/18$

How to calculate it? → BAYES THEOREM!

$$p(\text{Sick} | \text{Positive}) = \frac{p(\text{Sick}) \cdot p(\text{Positive} | \text{Sick})}{p(\text{Positive})}$$


# Generative Modelling

---

In generative modelling, we encode the way how we believe the observed data was generated in a probabilistic model.

# Generative Modelling: Example 1


## Cluster Model



# Generative Modelling: Example 1

## Cluster Model


- How many clusters? Easy: 2



# Generative Modelling: Example 1


## Cluster Model

- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$



# Generative Modelling: Example 1

## Cluster Model




- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$
- Cluster centers,  $k = 0, 1$ :  
 $\mu_x^k \sim \text{Uniform}(-3,3)$ ,  $\mu_y^k \sim \text{Uniform}(-3,3)$



# Generative Modelling: Example 1


## Cluster Model



- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$
- Cluster centers,  $k = 0, 1$ :  
 $\mu_x^k \sim \text{Uniform}(-3,3)$ ,  $\mu_y^k \sim \text{Uniform}(-3,3)$
- Cluster spread,  $k = 0, 1$ :  
 $\sigma^k \sim \text{InverseGamma}(1,1)$

# Generative Modelling: Example 1

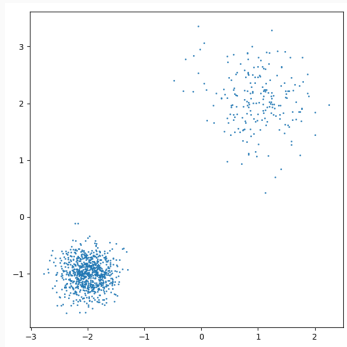
## Cluster Model



- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$
- Cluster centers,  $k = 0, 1$ :  
 $\mu_x^k \sim \text{Uniform}(-3,3)$ ,  $\mu_y^k \sim \text{Uniform}(-3,3)$
- Cluster spread,  $k = 0, 1$ :  
 $\sigma^k \sim \text{InverseGamma}(1,1)$
- Cluster membership,  $i = 1, \dots, N$ :  
 $z_i \sim \text{Bernoulli}(p)$

# Generative Modelling: Example 1


## Cluster Model



- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$
- Cluster centers,  $k = 0, 1$ :  
 $\mu_x^k \sim \text{Uniform}(-3,3)$ ,  $\mu_y^k \sim \text{Uniform}(-3,3)$
- Cluster spread,  $k = 0, 1$ :  
 $\sigma^k \sim \text{InverseGamma}(1,1)$
- Cluster membership,  $i = 1, \dots, N$ :  
 $z_i \sim \text{Bernoulli}(p)$
- Observed data,  $i = 1, \dots, N$ :  
 $x_i \sim \text{Normal}(\mu^{z_i}, \sigma^{z_i})$

# Generative Modelling: Example 1

## Cluster Model



- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$
- Cluster centers,  $k = 0, 1$ :  
 $\mu_x^k \sim \text{Uniform}(-3,3)$ ,  $\mu_y^k \sim \text{Uniform}(-3,3)$
- Cluster spread,  $k = 0, 1$ :  
 $\sigma^k \sim \text{InverseGamma}(1,1)$
- Cluster membership,  $i = 1, \dots, N$ :  
 $z_i \sim \text{Bernoulli}(p)$
- Observed data,  $i = 1, \dots, N$ :  
 $x_i \sim \text{Normal}(\mu^{z_i}, \sigma^{z_i})$


$$P(X, \Theta) = P(p)P(\mu_x^1)P(\mu_y^1)P(\mu_x^2)P(\mu_y^2)P(\sigma^1)P(\sigma^2) \prod_{i=1}^N P(z_i|p)P(x_i|\mu, \sigma, z_i)$$

# Generative Modelling: Example 1

- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$
- Cluster centers,  $k = 0, 1$ :  
 $\mu_x^k \sim \text{Uniform}(-3,3)$ ,  $\mu_y^k \sim \text{Uniform}(-3,3)$
- Cluster spread,  $k = 0, 1$ :  
 $\sigma^k \sim \text{InverseGamma}(1,1)$
- Cluster membership,  $i = 1, \dots, N$ :  
 $z_i \sim \text{Bernoulli}(p)$
- Observed data,  $i = 1, \dots, N$ :  
 $x_i \sim \text{Normal}(\mu^{z_i}, \sigma^{z_i})$

```
@model function cluster(x::Vector{Vector{Float64}})
    K = 2 # 2 clusters
    p ~ Uniform(0.,1) # probability of being in cluster 1
    mu_x = zeros(K)
    mu_y = zeros(K)
    sigmas = zeros(K)
    for k in 1:K
        # Cluster centers
        mu_x[k] ~ Uniform(-3,3)
        mu_y[k] ~ Uniform(-3,3)
        # Cluster spread
        sigmas[k] ~ InverseGamma(1,1)
    end
    z = zeros{Int, length(x)}
    for i in 1:length(x)
        z[i] ~ Bernoulli(p) # Cluster membership
        mu = z[i] == 1 ? [mu_x[1],mu_y[1]] : [mu_x[2],mu_y[2]]
        sigma = z[i] == 1 ? sigmas[1] : sigmas[2]
        x[i] ~ MvNormal(mu, sigma) # Observed data
    end
end
```

## Generative Modelling: Example 2



# Probabilistic Programming Languages

---

*Probabilistic programming is a programming paradigm in which probabilistic models are specified and posterior inference for these models is performed automatically.*



In the most general sense:

A **probabilistic program** is a non-deterministic program for which each execution can be weighted with a real number (e.g a probability).

# Probabilistic Programming Languages

```
1 def linear_regression(x, y):
2     lp = torch.tensor(0.0)
3
4     slope = dist.Normal(0., 3.).sample()
5     lp += dist.Normal(0., 3.).log_prob(slope)
6
7     intercept = dist.Normal(0., 3.).sample()
8     lp += dist.Normal(0., 3.).log_prob(intercept)
9
10    sigma = dist.HalfCauchy(1).sample()
11    lp += dist.HalfCauchy(1).log_prob(sigma)
12
13    for i in range(len(x)):
14        lp += dist.Normal(slope * x[i] + intercept, sigma).log_prob(y[i])
15
16
17    return (slope, intercept, sigma), lp
```

# Probabilistic Programming Languages

```
1 def linear_regression(x, y):
2     lp = torch.tensor(0.0)
3
4     slope = dist.Normal(0., 3.).sample()
5     lp += dist.Normal(0., 3.).log_prob(slope)
6
7     intercept = dist.Normal(0., 3.).sample()
8     lp += dist.Normal(0., 3.).log_prob(intercept)
9
10    sigma = dist.HalfCauchy(1).sample()
11    lp += dist.HalfCauchy(1).log_prob(sigma)
12
13    for i in range(len(x)):
14        lp += dist.Normal(slope * x[i] + intercept, sigma).log_prob(y[i])
15
16    return (slope, intercept, sigma), lp
```

$$P(X, \Theta) = P(\text{slope})P(\text{intercept})P(\text{sigma}) \prod_{i=1}^N P(x_i, y_i | \text{slope}, \text{intercept}, \text{sigma})$$

# Probabilistic Programming Languages

```
1 def linear_regression_lp(slope, intercept, sigma, x, y):
2     lp = torch.tensor(0.0)
3
4     lp += dist.Normal(0., 3.).log_prob(slope)
5
6     lp += dist.Normal(0., 3.).log_prob(intercept)
7
8     lp += dist.HalfCauchy(1).log_prob(sigma)
9
10    for i in range(len(x)):
11        lp += dist.Normal(slope * x[i] + intercept, sigma).log_prob(y[i])
12
13    return lp
```

A **probabilistic programming language (PPL)** abstracts away the complexities of computing the program weight.

# Probabilistic Programming Languages

A **probabilistic programming language (PPL)** abstracts away the complexities of computing the program weight.

A PPL allows us to conveniently specify a probabilistic model as a program (by using special syntax or primitives).

# Probabilistic Programming Languages

A **probabilistic programming language (PPL)** abstracts away the complexities of computing the program weight.

A PPL allows us to conveniently specify a probabilistic model as a program (by using special syntax or primitives).

Internally, the programs are represented in a way to facilitate automated posterior inference.

# Probabilistic Programming Languages

```
1 # Turing
2 @model function linreg(x, y)
3     slope ~ Normal(0,3)
4     intercept ~ Normal(0,3)
5     sigma ~ InverseGamma(1,1)
6     for i in eachindex(y)
7         y[i] ~ Normal(slope * x[i] + intercept, sigma)
8     end
9 end
```



# Probabilistic Programming Languages

```
1 # Turing
2 @model function linreg(x, y)
3     slope ~ Normal(0,3)
4     intercept ~ Normal(0,3)
5     sigma ~ InverseGamma(1,1)
6     for i in eachindex(y)
7         y[i] ~ Normal(slope * x[i] + intercept, sigma)
8     end
9 end

1 # Pyro
2 def linreg(x, y):
3     slope = pyro.sample("slope", dist.Normal(0,3))
4     intercept = pyro.sample("intercept", dist.Normal(0,3))
5     sigma = pyro.sample("sigma", dist.HalfCauchy(1))
6
7     for i in range(len(x)):
8         pyro.sample(
9             f"y[{i}]",
10             dist.Normal(slope * x[i] + intercept, sigma),
11             obs=y[i])
```

## Representation-Based

- Translate the program to internal representation
- Usually the underlying representation is some sort of graph
- Optimised inference

## Representation-Based

- Translate the program to internal representation
- Usually the underlying representation is some sort of graph
- Optimised inference

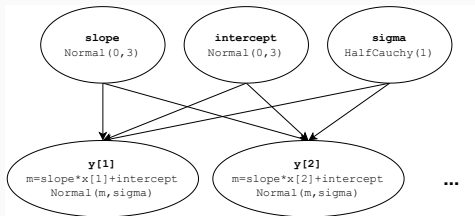
## Evaluation-Based

- Usually embedded in existing languages
- Have special programming constructs to specify models
- Run the entire program in different contexts for inference
- General Purpose

# Representation-Based: Graphical Models

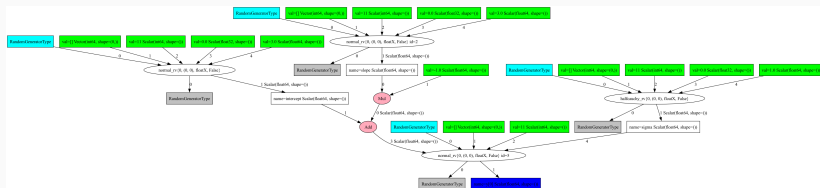
## Example: PyMC

```
1 x = np.array([-1.0, -0.5, 0.0, 0.5, 1.0])
2 y = np.array([-1.2, -1.5, -0.0, -0.8, 1.5])
3 with pymc.Model() as linreg:
4     slope = pymc.Normal("slope", 0.,3.)
5     intercept = pymc.Normal("intercept", 0.,3.)
6     sigma = pymc.HalfCauchy("sigma", 1)
7     for i in range(len(y)):
8         pymc.Normal(
9             f"y[{i}]",
10             slope * x[i] + intercept, sigma,
11             observed=y[i])
```



# Representation-Based: Graphical Models

## Underlying Computation Graph in PyMC



## Strengths

- Can exploit structure of model
- Optimised inference

## Constraints

- Only works for models with finite number of random variables
- Are not "embedded" in host language

# Representation-Based: Unconstrained Function

## Example: Stan

```
data {  
  int<lower=0> N;  
  vector[N] x;  
  vector[N] y;  
}  
parameters {  
  real slope;  
  real intercept;  
  real<lower=0> sigma;  
}  
model {  
  slope ~ normal(0,3)  
  intercept ~ normal(0,3)  
  sigma ~ inv_gamma(2,1)  
  y ~ normal(slope+intercept*x,sigma);  
}  
  
double model(vector<double> theta,  
  int N, vector<double> x, vector<double> y) {  
  
  double target = 0.0;  
  
  slope = theta[0];  
  intercept = theta[1];  
  sigma = exp(theta[2]); // transformed  
  
  target += normal_lpdf(slope, 0, 3);  
  target += normal_lpdf(intercept, 0, 3);  
  target += inv_gamma_lpdf(sigma, 2, 1);  
  
  for (int i = 0; i++, i < N) {  
    target += normal_lpdf(y[i], slope*x[i]+intercept, sigma)  
  }  
  
  return target;  
}
```

$$\text{model: } \mathbb{R}^3 \rightarrow \mathbb{R}$$

a.s. differentiable if all variables are continuous and terminates

# Representation-Based: Unconstrained Function

## Strengths

- Specialised (gradient-based) inference
- Fast inference

## Constraints

- Black box function
- Only supports finite number of random variables



## Example: Psi

```
def main(){  
  flips := [0,1,1,0,0];  
  n := flips.length;  
  
  p := uniform(0,1);  
  
  for i in [0..n) {  
    x := bernoulli(p);  
    observe(x==flips[i]);  
  }  
  
  u := bernoulli(p);  
  return u;  
}
```

$$p(u) = \frac{\int_0^1 p^{2+u}(1-p)^{4-u} dp}{\int_0^1 p^3(1-p)^3 dp + \int_0^1 p^2(1-p)^4 dp}$$

## Example: Psi

```
def main(){  
  flips := [0,1,1,0,0];  
  n := flips.length;  
  
  p := uniform(0,1);  
  
  for i in [0..n) {  
    x := bernoulli(p);  
    observe(x==flips[i]);  
  }  
  
  u := bernoulli(p);  
  return u;  
}
```

$$\begin{aligned} p(u) &= \frac{\int_0^1 p^{2+u}(1-p)^{4-u} dp}{\int_0^1 p^3(1-p)^3 dp + \int_0^1 p^2(1-p)^4 dp} \\ &= \frac{\int_0^1 p^{2+u}(1-p)^{4-u} dp}{\frac{1}{140} + \frac{1}{105}} \\ &= \frac{3}{7} \delta_1(u) + \frac{4}{7} \delta_0(u) \end{aligned}$$

## Strengths

- Exact inference

## Constraints

- Most models are not solvable symbolically

# Evaluation-Based

## Example: Pyro

```
x = torch.tensor([-1.0, -0.5, 0.0, 0.5, 1.0])
y = torch.tensor([-1.2, -1.5, -0.0, -0.8, 1.5])
```

```
def linreg(x):
    slope = pyro.sample("slope", dist.Normal(0,3))
    intercept = pyro.sample("intercept", dist.Normal(0,3))
    sigma = pyro.sample("sigma", dist.HalfCauchy(1))

    for i in range(len(x)):
        pyro.sample(
            f"y[{i}]",
            dist.Normal(slope * x[i] + intercept, sigma))

    return slope, intercept
```

```
linreg(x)
returns:
(tensor(0.4707), tensor(3.9712))
```

```
with pyro.poutine.trace() as trace:
    linreg(x)
trace.get_trace().nodes["sigma"]
returns:
{'type': 'sample',
 'name': 'sigma',
 'fn': HalfCauchy(),
 'is_observed': False,
 'value': tensor(4.4170),
 ...}
```

# Evaluation-Based

## Example: Pyro

```
x = torch.tensor([-1.0, -0.5, 0.0, 0.5, 1.0])
y = torch.tensor([-1.2, -1.5, -0.0, -0.8, 1.5])

def linreg(x):
    slope = pyro.sample("slope", dist.Normal(0,3))
    intercept = pyro.sample("intercept", dist.Normal(0,3))
    sigma = pyro.sample("sigma", dist.HalfCauchy(1))

    for i in range(len(x)):
        pyro.sample(
            f"y[{i}]",
            dist.Normal(slope * x[i] + intercept, sigma))

    return slope, intercept
```

```
with pyro.poutine.trace() as trace:
    with pyro.condition(
        data={f"y[{i}]": y[i] for i in range(5)}):
        linreg(x)
trace.get_trace().nodes["y[0]"]
returns:
{'type': 'sample',
 'name': 'y[0]',
 'fn': Normal(loc: -7.59, scale: 2.31),
 'is_observed': True,
 'value': tensor(-1.2000),
 ...}

trace.get_trace().log_prob_sum()
returns:
tensor(-27.8861)
```

## Example: Gen

```
@gen function linreg(x, y)
  slope ~ normal(0,3)
end
```

```
function var"##linreg#369"(var"##state#294", x::Any, y::Any)
  slope = Gen.traceat(var"##state#294", normal, (0, 3), :slope)
end
```

# Evaluation-Based

```
@model function linreg(x, y) # Turing
  slope ~ Normal(0,3)
end
```

```
function linreg(__model__::Model, __varinfo__::AbstractVarInfo, __context__::AbstractContext, x, y)
  var"##dist#362" = Normal(0, 3)
  var"##vn#359" = (DynamicPPL.resolve_varnames)((AbstractPPL.VarName){:slope}(), var"##dist#362")
  var"##isassumption#360" = begin
    if (DynamicPPL.contextual_isassumption)(__context__, var"##vn#359")
      ...
    end
  end
  begin
    var"##retval#364" = if (DynamicPPL.contextual_isfixed)(__context__, var"##vn#359")
      slope = (DynamicPPL.getfixed_nested)(__context__, var"##vn#359")
      elseif var"##isassumption#360"
        begin
          (var"##value#363", __varinfo__) = (DynamicPPL.tilde_assume!!)(__context__, (DynamicPPL.unwr
            slope = var"##value#363"
            var"##value#363"
          end
        end
      else
        if !((DynamicPPL.inargnames)(var"##vn#359", __model__))
          slope = (DynamicPPL.getconditioned_nested)(__context__, var"##vn#359")
        end
        (var"##value#361", __varinfo__) = (DynamicPPL.tilde_observe!!)(__context__, (DynamicPPL.check_ti
          var"##value#361"
        end
      end
    return (var"##retval#364", __varinfo__)
  end
end
```

## Strengths

- Can represent any probabilistic model
- Can be embedded in host languages

## Constraints

- General-purpose inference is difficult



## Implementing a minimal PPL

---

# Implementing a minimal PPL

## Goals:

- Evaluation-based
- Contextualised execution (*Lecture 2*)

## Goals:

- Evaluation-based
- Contextualised execution (*Lecture 2*)
- General-purpose inference:
  - Likelihood weighting (*Lecture 2 + Assignment 2*)
  - Metropolis Hastings (*Lecture 3 + Assignment 3*)

# Implementing a minimal PPL

## Goals:

- Evaluation-based
- Contextualised execution (*Lecture 2*)
- General-purpose inference:
  - Likelihood weighting (*Lecture 2 + Assignment 2*)
  - Metropolis Hastings (*Lecture 3 + Assignment 3*)
- Gradient-based inference:
  - Hamiltonian Monte Carlo (*Lecture 3 + Assignment 4*)
  - Variational Inference (*Lecture 4 + Assignment 4*)

## torch.distributions

```
1 class Distribution:
2     def sample(self, sample_shape: torch.Size = torch.Size()) -> torch.Tensor:
3         raise NotImplementedError
4
5     def log_prob(self, value: torch.Tensor) -> torch.Tensor:
6         raise NotImplementedError
```

# PPL Core: Sample Context

```
1  _SAMPLE_CONTEXT = None
2
3  class SampleContext(ABC):
4      def __enter__(self):
5          global _SAMPLE_CONTEXT
6          _SAMPLE_CONTEXT = self
7
8      def __exit__(self, *args):
9          global _SAMPLE_CONTEXT
10         _SAMPLE_CONTEXT = None
11
12     @abstractmethod
13     def sample(self,
14               address: str,
15               distribution: dist.Distribution,
16               observed: Optional[torch.Tensor] = None) -> torch.Tensor:
17         raise NotImplementedError
```

# PPL Core: Sample Context

```
1 ctx = MySampleContext()
2 with ctx:
3     do_something()
```

```
1 ctx = MySampleContext()
2 ctx.__enter__() # global _SAMPLE_CONTEXT = ctx
3 do_something()
4 ctx.__exit__() # global _SAMPLE_CONTEXT = None
```

# PPL Core: Sample Statement

```
1 def sample(address: str,  
2           distribution: dist.Distribution,  
3           observed: Optional[torch.Tensor] = None) -> torch.Tensor:  
4     global _SAMPLE_CONTEXT  
5  
6     # default behavior  
7     if _SAMPLE_CONTEXT is None:  
8         if observed is not None:  
9             return observed  
10        return distribution.sample()  
11  
12    # context specific behavior  
13    return _SAMPLE_CONTEXT.sample(address, distribution, observed)
```



# Example Program

```
1 def noisy_geometric(p):
2     x = 0
3     while True:
4         b = sample(f"b_{x}", dist.Bernoulli(p))
5         if b:
6             break
7         x += 1
8     y = sample("y", dist.Normal(x,1), observed=torch.tensor(3.))
9     return x
10
```

```
torch.manual_seed(0)
[noisy_geometric(0.25) for _ in range(10)]
```

**returns:**

```
[2, 0, 8, 0, 4, 8, 0, 3, 1, 0]
```

# Example Context

```
1  class LogProb(SampleContext):
2      def __init__(self, trace):
3          self.log_prob = torch.tensor(0.)
4          self.trace = trace
5
6      def sample(self,
7                 address: str,
8                 distribution: dist.Distribution,
9                 observed: Optional[torch.Tensor] = None) -> torch.Tensor:
10
11         if observed is not None:
12             value = observed
13         else:
14             value = self.trace[address]
15
16         self.log_prob += distribution.log_prob(value)
17
18         return value
19
```

## Example Context (cont.)

```
torch.manual_seed(0)
ctx = LogProb({
    "b_0": torch.tensor(0.),
    "b_1": torch.tensor(0.),
    "b_2": torch.tensor(1.)})
with ctx:
    x = noisy_geometric(0.25)
x, ctx.log_prob
```

**returns:**

```
(2, tensor(-3.3806))
```

```
p = torch.tensor(0.25)
2*torch.log(1-p) + torch.log(p) + dist.Normal(2,1).log_prob(torch.tensor(3))
# == tensor(-3.3806)
```

# Example Context

```
1  class Trace(SampleContext):
2      def __init__(self):
3          self.trace = {}
4
5      def sample(self,
6                  address: str,
7                  distribution: dist.Distribution,
8                  observed: Optional[torch.Tensor] = None) -> torch.Tensor:
9
10         if observed is not None:
11             value = observed
12         else:
13             value = distribution.sample()
14
15         self.trace[address] = {
16             'value': value,
17             'distribution': distribution,
18             'is_observed': observed is not None,
19             'log_prob': distribution.log_prob(value)
20         }
21
22         return value
23
```

## Example Context (cont.)


```
torch.manual_seed(0)
ctx = Trace()
with ctx:
    x = noisy_geometric(0.25)
ctx.trace
```

returns:

```
{'b_0': {'value': tensor(0.),
        'distribution': Bernoulli(probs: 0.25, logits: -1.0986123085021973),
        'is_observed': False,
        'log_prob': tensor(-0.2877)},
 'b_1': {'value': tensor(0.),
        'distribution': Bernoulli(probs: 0.25, logits: -1.0986123085021973),
        'is_observed': False,
        'log_prob': tensor(-0.2877)},
 'b_2': {'value': tensor(1.),
        'distribution': Bernoulli(probs: 0.25, logits: -1.0986123085021973),
        'is_observed': False,
        'log_prob': tensor(-1.3863)},
 'y': {'value': tensor(3),
       'distribution': Normal(loc: 2.0, scale: 1.0),
       'is_observed': True,
       'log_prob': tensor(-1.4189)}}
```

# Generative Modelling: Example 1

## Cluster Model



- How many clusters? Easy: 2
- Model probability of being in cluster 1:  
 $p \sim \text{Uniform}(0,1)$
- Cluster centers,  $k = 0, 1$ :  
 $\mu_x^k \sim \text{Uniform}(-3,3)$ ,  $\mu_y^k \sim \text{Uniform}(-3,3)$
- Cluster spread,  $k = 0, 1$ :  
 $\sigma^k \sim \text{InverseGamma}(1,1)$
- Cluster membership,  $i = 1, \dots, N$ :  
 $z_i \sim \text{Bernoulli}(p)$
- Observed data,  $i = 1, \dots, N$ :  
 $x_i \sim \text{Normal}(\mu^{z_i}, \sigma^{z_i})$

$$P(X, \Theta) = P(p)P(\mu_x^1)P(\mu_y^1)P(\mu_x^2)P(\mu_y^2)P(\sigma^1)P(\sigma^2) \prod_{i=1}^N P(z_i|p)P(x_i|\mu, \sigma, z_i)$$

# Example Program

```
def cluster(x):
    K = 2 # 2 clusters
    p = sample("p", dist.Uniform(0.,1)) # probability of being in cluster 1
    mu_x = []
    mu_y = []
    sigmas = []
    for k in range(K):
        # Cluster centers
        mu_x.append(sample(f"mu_x_{k}", dist.Uniform(-3,3)))
        mu_y.append(sample(f"mu_y_{k}", dist.Uniform(-3,3)))
        # Cluster spread
        sigmas.append(sample(f"sigma_{k}", dist.HalfCauchy(1)))
    for i in range(len(x)):
        z = sample(f"z_{i}", dist.Bernoulli(p)) # Cluster membership
        mu = torch.hstack([mu_x[0],mu_y[0]]) if z == 1 else torch.hstack([mu_x[1],mu_y[1]])
        sigma = sigmas[0] if z == 1 else sigmas[1]
        sample(f"x_{i}", dist.Normal(mu, sigma), observed=x[i]) # Observed data
```

# First General-Purpose Inference

---



# Rejection Sampling

Coin model:

$$p \sim \text{Uniform}(0, 1)$$

$$x_i \sim \text{Bernoulli}(p)$$

```
1 def rejection_sampling(xs_observed):
2     while True:
3         # generate samples from joint
4         p = dist.Uniform(0, 1).sample()
5         xs = dist.Bernoulli(p).sample(xs_observed.shape)
6         if (xs == xs_observed).all():
7             return p # accept
8         # reject
```

# Rejection Sampling

Coin model:

$$p \sim \text{Uniform}(0, 1)$$

$$x_i \sim \text{Bernoulli}(p)$$

```
1 def rejection_sampling(xs_observed):
2     while True:
3         # generate samples from joint
4         p = dist.Uniform(0, 1).sample()
5         xs = dist.Bernoulli(p).sample(xs_observed.shape)
6         if (xs == xs_observed).all():
7             return p # accept
8         # reject
```

*Rejected 50.54% of iterations for 1 number of observations.  
Rejected 83.78% of iterations for 2 number of observations.  
Rejected 91.76% of iterations for 3 number of observations.  
Rejected 96.68% of iterations for 4 number of observations.  
Rejected 98.26% of iterations for 5 number of observations.*

## Likelihood Weighting (Importance Sampling)

Problem: We want to compute statistic  $r(\Theta)$  for the posterior  $P(\Theta|X)$ , but we cannot sample from it or evaluate the density directly.

Solution: Sample from a reference distribution  $Q$ .

## Likelihood Weighting (Importance Sampling)

Problem: We want to compute statistic  $r(\Theta)$  for the posterior  $P(\Theta|X)$ , but we cannot sample from it or evaluate the density directly.

Solution: Sample from a reference distribution  $Q$ .

$$\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)]$$

Posterior mean:  $r(\Theta) = \Theta$ , probability that  $\Theta > 0$ :  $r(\Theta) = \delta_{\Theta > 0}$ .

## Importance Sampling

$$\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)]$$

## Importance Sampling

$$\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] = \int r(\theta) p_{\Theta|X}(\theta|X) d\theta$$

## Importance Sampling

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \int r(\theta) p_{\Theta|X}(\theta|X) d\theta \\ &= \int r(\theta) \frac{p_{\Theta|X}(\theta|X)}{p_Q(\theta)} p_Q(\theta) d\theta\end{aligned}$$

# First General-Purpose Inference

## Importance Sampling

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \int r(\theta) p_{\Theta|X}(\theta|X) d\theta \\ &= \int r(\theta) \frac{p_{\Theta|X}(\theta|X)}{p_Q(\theta)} p_Q(\theta) d\theta \\ &= \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right]\end{aligned}$$



# First General-Purpose Inference

## Importance Sampling

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \int r(\theta) p_{\Theta|X}(\theta|X) d\theta \\ &= \int r(\theta) \frac{p_{\Theta|X}(\theta|X)}{p_Q(\theta)} p_Q(\theta) d\theta \\ &= \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right]\end{aligned}$$

We can use samples  $\theta_i$  from  $Q$  in Monte Carlo simulation:

$$\mathbb{E}_{X \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \approx \frac{1}{N} \sum_{i=1}^N r(\theta_i) \frac{p_{\Theta|X}(\theta_i|X)}{p_Q(\theta_i)}$$

# First General-Purpose Inference

## Importance Sampling

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \int r(\theta) p_{\Theta|X}(\theta|X) d\theta \\ &= \int r(\theta) \frac{p_{\Theta|X}(\theta|X)}{p_Q(\theta)} p_Q(\theta) d\theta \\ &= \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right]\end{aligned}$$

We can use samples  $\theta_i$  from  $Q$  in Monte Carlo simulation:

$$\mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \approx \frac{1}{N} \sum_{i=1}^N r(\theta_i) \frac{p_{\Theta|X}(\theta_i|X)}{p_Q(\theta_i)}$$

But we cannot compute  $p_{\Theta|X}(\theta_i|X)$ !?

But we cannot compute  $p_{\Theta|X}(\theta_i|X)$ !?

$$\frac{p_{\Theta|X}(\theta_i|X)}{p_Q(\theta_i)}$$

But we cannot compute  $p_{\Theta|X}(\theta_i|X)$ !?

$$\frac{p_{\Theta|X}(\theta_i|X)}{p_Q(\theta_i)} = \frac{p_{\Theta,X}(\theta_i, X)}{p_X(X)p_Q(\theta_i)}$$

# First General-Purpose Inference

But we cannot compute  $p_{\Theta|X}(\theta_i|X)$ !?

$$\begin{aligned}\frac{p_{\Theta|X}(\theta_i|X)}{p_Q(\theta_i)} &= \frac{p_{\Theta,X}(\theta_i, X)}{p_X(X)p_Q(\theta_i)} \\ &= \frac{1}{p_X(X)} \underbrace{\frac{p_{\Theta,X}(\theta_i, X)}{p_Q(\theta_i)}}_{W_i :=}\end{aligned}$$

But we cannot compute  $p_X(X)$ !?

# First General-Purpose Inference

But we cannot compute  $p_X(X)$ !?

Set  $r(\Theta) = 1$

$$1 = \mathbb{E}_{X \sim P(\cdot|X)} [r(\Theta)]$$

# First General-Purpose Inference

But we cannot compute  $p_X(X)$ !?

Set  $r(\Theta) = 1$

$$1 = \mathbb{E}_{X \sim P(\cdot|X)} [r(\Theta)] = \mathbb{E}_{\Theta \sim Q} \left[ \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right]$$

# First General-Purpose Inference

But we cannot compute  $p_X(X)$ !?

Set  $r(\Theta) = 1$

$$\begin{aligned} 1 = \mathbb{E}_{X \sim P(\cdot|X)} [r(\Theta)] &= \mathbb{E}_{\Theta \sim Q} \left[ \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \\ &= \frac{1}{p_X(X)} \mathbb{E}_{\Theta \sim Q} \left[ \frac{p_{\Theta,X}(\Theta, X)}{p_Q(\Theta)} \right] \end{aligned}$$



# First General-Purpose Inference

But we cannot compute  $p_X(X)$ !?

Set  $r(\Theta) = 1$

$$\begin{aligned} 1 = \mathbb{E}_{X \sim P(\cdot|X)} [r(\Theta)] &= \mathbb{E}_{\Theta \sim Q} \left[ \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \\ &= \frac{1}{p_X(X)} \mathbb{E}_{\Theta \sim Q} \left[ \frac{p_{\Theta,X}(\Theta, X)}{p_Q(\Theta)} \right] \approx \frac{1}{p_X(X)} \frac{1}{N} \sum_{i=1}^N W_i \end{aligned}$$

# First General-Purpose Inference

But we cannot compute  $p_X(X)$ !?

Set  $r(\Theta) = 1$

$$\begin{aligned} 1 = \mathbb{E}_{X \sim P(\cdot|X)} [r(\Theta)] &= \mathbb{E}_{\Theta \sim Q} \left[ \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \\ &= \frac{1}{p_X(X)} \mathbb{E}_{\Theta \sim Q} \left[ \frac{p_{\Theta,X}(\Theta, X)}{p_Q(\Theta)} \right] \approx \frac{1}{p_X(X)} \frac{1}{N} \sum_{i=1}^N W_i \end{aligned}$$

Thus,

$$p_X(X) \approx \frac{1}{N} \sum_{i=1}^N W_i$$

# First General-Purpose Inference

Putting it all together:  $W_i = \frac{p_{\Theta, X}(\theta_i, X)}{p_Q(\theta_i)}$

$$\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)]$$

# First General-Purpose Inference

Putting it all together:  $W_i = \frac{p_{\Theta, X}(\theta_i, X)}{p_Q(\theta_i)}$

$$\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] = \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right]$$

# First General-Purpose Inference

Putting it all together:  $W_i = \frac{p_{\Theta, X}(\theta_i, X)}{p_Q(\theta_i)}$

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \\ &\approx \frac{1}{N} \sum_{i=1}^N \frac{p_{\Theta, X}(\theta_i, X)}{p_X(X)p_Q(\theta_i)} \cdot r(\theta_i)\end{aligned}$$

# First General-Purpose Inference

Putting it all together:  $W_i = \frac{p_{\Theta, X}(\theta_i, X)}{p_Q(\theta_i)}$

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \\ &\approx \frac{1}{N} \sum_{i=1}^N \frac{p_{\Theta, X}(\theta_i, X)}{p_X(X) p_Q(\theta_i)} \cdot r(\theta_i) \\ &= \frac{1}{N} \sum_{i=1}^N \frac{W_i}{p_X(X)} r(\theta_i)\end{aligned}$$

# First General-Purpose Inference

Putting it all together:  $W_i = \frac{p_{\Theta, X}(\theta_i, X)}{p_Q(\theta_i)}$

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \\ &\approx \frac{1}{N} \sum_{i=1}^N \frac{p_{\Theta, X}(\theta_i, X)}{p_X(X)p_Q(\theta_i)} \cdot r(\theta_i) \\ &= \frac{1}{N} \sum_{i=1}^N \frac{W_i}{p_X(X)} r(\theta_i) \\ &\approx \frac{1}{N} \sum_{i=1}^N \frac{W_i}{\sum_{k=1}^N W_k} r(\theta_i), \quad \theta_i \sim Q\end{aligned}$$

# First General-Purpose Inference

Putting it all together:  $W_i = \frac{p_{\Theta, X}(\theta_i, X)}{p_Q(\theta_i)}$

$$\begin{aligned}\mathbb{E}_{\Theta \sim P(\cdot|X)} [r(\Theta)] &= \mathbb{E}_{\Theta \sim Q} \left[ r(\Theta) \frac{p_{\Theta|X}(\Theta|X)}{p_Q(\Theta)} \right] \\ &\approx \frac{1}{N} \sum_{i=1}^N \frac{p_{\Theta, X}(\theta_i, X)}{p_X(X)p_Q(\theta_i)} \cdot r(\theta_i) \\ &= \frac{1}{N} \sum_{i=1}^N \frac{W_i}{p_X(X)} r(\theta_i) \\ &\approx \frac{1}{N} \sum_{i=1}^N \frac{W_i}{\sum_{k=1}^N W_k} r(\theta_i), \quad \theta_i \sim Q\end{aligned}$$

$Q$  has to satisfy:  $p_Q(\theta) = 0 \implies p_{\Theta, X}(\theta, X) = 0$ .



**Likelihood Weighting** is a special case of importance sampling where the reference distribution is the prior  $Q = P_{\Theta}$ .

The weights reduce to the likelihood

$$w_i = \frac{p_{\Theta, X}(\theta_i, X)}{p_Q(\theta_i)} = \frac{p_{\Theta, X}(\theta_i, X)}{p_{\Theta}(\theta_i)} = p_{X|\Theta}(X|\theta_i)$$


# Example

## Coin model

```
1 def prior(p):
2     return dist.Uniform(0,1, validate_args=False).log_prob(p).exp()
3
4 def likelihood(xs,p):
5     lp = torch.tensor(0.)
6     for x in xs:
7         lp += dist.Bernoulli(p, validate_args=False).log_prob(x)
8     return lp.exp()
9
10 def joint(xs, p):
11     return prior(p) * likelihood(xs, p)
12
13 xs = torch.tensor([0.,1.,1.,0.,0.])
14
```


# Example

```
1 proposal = dist.Normal(0,1)
2 proposal_sample = proposal.sample((N,))
3
4 weights = (
5     torch.hstack([joint(xs, proposal_sample[i]) for i in range(N)]) /
6     proposal.log_prob(proposal_sample).exp()
7 ) # importance sampling
8
9 axs[0].hist(proposal_sample, density=True, bins=100)
10 axs[1].hist(proposal_sample, density=True, bins=100, weights=weights)
```



# Example

```
1 proposal = dist.Uniform(0,1) # prior
2 proposal_sample = proposal.sample((N,))
3
4 weights = torch.hstack([likelihood(xs, proposal_sample[i]) for i in range(N)])
5 # likelihood weighting
6
7 axs[0].hist(proposal_sample, density=True, bins=100)
8 axs[1].hist(proposal_sample, density=True, bins=100, weights=weights)
9
```



# Example

In our PPL, the coin model can be specified much more compactly.

```
1 def coin_model(xs):  
2     p = sample("p", dist.Uniform(0,1))  
3     for i in range(len(xs)):  
4         sample(f"x[{i}]", dist.Bernoulli(p), observed=xs[i])
```

In assignment 2, we will write a sample context to make likelihood weighting work automatically with any model.

# Resources

Bayesian Inference Framework

[\*https://www.youtube.com/watch?v=0w\\_4QcvBYII\*](https://www.youtube.com/watch?v=0w_4QcvBYII)

Intuition behind Bayesian inference

[\*https://www.youtube.com/watch?v=yvWlpwnT1nw\*](https://www.youtube.com/watch?v=yvWlpwnT1nw)

Bayesian posterior sampling

[\*https://www.youtube.com/watch?v=EHqU9LE9tg8\*](https://www.youtube.com/watch?v=EHqU9LE9tg8)

Generative Models

[\*http://probmods.org/chapters/generative-models.html\*](http://probmods.org/chapters/generative-models.html)

AI That Understands the World, Using Probabilistic Programming

[\*https://www.youtube.com/watch?v=8j2S7BRRWus\*](https://www.youtube.com/watch?v=8j2S7BRRWus)

A Personal Viewpoint on Probabilistic Programming

[\*https://www.youtube.com/watch?v=TFXcVlKqPlM\*](https://www.youtube.com/watch?v=TFXcVlKqPlM)